

Object-Oriented Programming and Data Structures

COMP2012: Sorting

Sorting Algorithms

- Problem: Given a list of N items, sort them in non-descending order.
- Comparison-based, exchanging adjacent elements
 - Bubble sort
 - Selection sort
 - Insertion sort
- Comparison-based, divide-and-conquer
 - Merge sort
 - Quicksort
- Non-comparison-based, linear time
 - Radix sort

Some Terminologies

- **Stable** algorithms:
 - do **not** change the **relative order** of elements with **equal** values.
 - Importance: e.g., sort names first on the first name and then on the last name.
 - if the last names are sorted by stable algorithm, names with the same last name will maintain the relative sorted order of their first name.
- **In-place** algorithms:
 - no extra space is required, or require only a constant amount $O(1)$ of extra memory space.
- **Online** algorithms:
 - sort a list as it receives it.

Comparison-based Sorting Algorithms by Exchanging Adjacent Elements

Bubble Sort

- **Procedure**
 1. Starting from the first item, compare each pair of **adjacent items**.
 2. If they are out of order, swap them.
 3. Repeat steps 1 – 2 with only the first $(N-j)$ items until **no** swapping occurs in an iteration, where j is the iteration index (starting from 1).
- **Idea:** At the end of each iteration, the maximum item should “bubble” to the end.
- **Worst-case:** $(N-1) + (N-2) + \dots + 1 = N(N+1)/2 \Rightarrow O(N^2)$
- **Average-case:** $O(N^2)$
- **Best-case:** $O(N)$ when the input is already sorted.

Bubble Sort: Example

- To sort: 2, 9, 7, 4, 1

- first pass:

- 2, 9, 7, 4, 1
 - 2, 7, 9, 4, 1
 - 2, 7, 4, 9, 1
 - 2, 7, 4, 1, 9

- second pass:

- 2, 7, 4, 1, 9
 - 2, 4, 7, 1, 9
 - 2, 4, 1, 7, 9

- etc.

Bubble Sort Code

```
#include <algorithm>

void bubble_sort(int* data, int N)
{
    bool swapped = true;

    while (swapped) // If swapping occurs in the previous iteration
    {
        swapped = false;

        for (int j = 1; j < N; ++j)
        {
            if (data[j] < data[j-1]) // Max item bubbles to the end
            {
                std::swap(data[j], data[j-1]);
                swapped = true;      // Swapping occurs
            }
        }
        --N;
    }
}
```

Selection Sort

- Procedure
 - 1. Find the **minimum** item in the list
 - 2. Swap it with the item in the **first** position
 - 3. Repeat the steps 1 and 2 for the rest of the list (starting at the second position)
- **Worst/Average/Best-case analysis:**
 - finding the minimum item is $O(N)$
 - the procedure is repeated N times => $O(N^2)$

Selection Sort: Example

- To sort: 2, 9, 7, 4, 1
 - 1st iteration: min_index = 4
 - 1, 9, 7, 4, 2
 - 2nd iteration: min_index = 3 for the sub-list
 - 1, 2, 7, 4, 9
 - 3rd iteration: min_index = 1 for the sub-list
 - 1, 2, 4, 7, 9
 - 4th iteration: min_index = 0 for the sub-list
 - 1, 2, 4, 7, 9
 - 5th iteration: return immediately
 - 1, 2, 4, 7, 9

Selection Sort Code

```
#include <algorithm>

void selection_sort(int* data, int N)
{
    if (N <= 1) // Base case
        return;

    // Find the index of the minimum value
    int min_index = 0;

    for (int j = 1; j < N; ++j)
    {
        if (data[min_index] > data[j])
            min_index = j;
    }

    if (min_index != 0)
        std::swap(data[0], data[min_index]); // Swap min with the first
    selection_sort(data+1, N-1); // Sort the remaining list
}
```

Insertion Sort

- Procedure

1. It takes $(N - 1)$ passes.
2. For pass $p = 1$ to $(N - 1)$, items in position 0 to $(p-1)$ are sorted.
3. In pass p , continuously move the item in position p left until it finds its **correct** position in the preceding p items.



Insertion Sort: Example

Original	34	8	64	51	32	21	Positions Moved
After $p = 1$	8	34	64	51	32	21	1
After $p = 2$	8	34	64	51	32	21	0
After $p = 3$	8	34	51	64	32	21	1
After $p = 4$	8	32	34	51	64	21	3
After $p = 5$	8	21	32	34	51	64	4

Insertion Sort Code

```
void insertion_sort(int* data, int N)
{
    for (int p = 1; p < N; ++p)
    {
        int item_to_insert = data[p];

        int j;
        for (j = p; j > 0 && item_to_insert < data[j-1]; --j)
            data[j] = data[j-1];

        data[j] = item_to_insert;
    }
}
```

Insertion Sort: Extended Example

- 6 items to sort: 34 8 64 51 32 21

`p = 1; item_to_insert = 8;`

`34 > item_to_insert`, so 2nd element `a[1]` is set to 34: {8, 34}...

We have reached the front of the list.

Thus, 1st position `a[0] = item_to_insert=8`

After 1st pass: 8 34 64 51 32 21

(first 2 elements are sorted)

`p = 2; item_to_insert = 64;`

`34 < 64`, so stop at 3rd position and set 3rd position = 64

After 2nd pass: 8 34 64 51 32 21

(first 3 elements are sorted)

Insertion Sort: Extended Example ..

p = 3; item_to_insert = 51;

51 < 64, so we have 8 34 64 64 32 21,

34 < 51, so stop at 2nd position, set 3rd position = item_to_insert,

After 3rd pass: 8 34 51 64 32 21

(first 4 elements are sorted)

p = 4; item_to_insert = 32;

32 < 64, so 8 34 51 64 64 21,

32 < 51, so 8 34 51 51 64 21,

next 32 < 34, so 8 34 34 51 64 21,

next 32 > 8, so stop at 1st position and set 2nd position = 32,

After 4th pass: 8 32 34 51 64 21

p = 5; item_to_insert = 21, ...

After 5th pass: 8 21 32 34 51 64

Insertion Sort: Worst-case Analysis

```
for (j = p; j > 0 && item_to_insert < data[j-1]; --j)
    data[j] = data[j-1];
```

- Inner loop is executed p times, for each $p = 1..N$
 \Rightarrow Overall: $1 + 2 + 3 + \dots + (N-1) = O(N^2)$
- Space (memory) requirement is $O(N)$

Insertion Sort: The Bound is Tight

- Tight bound: $\Theta(N^2)$
- That is, there exists some input which actually uses (the lower bound) $\Omega(N^2)$ time.
- Consider the worst case with a reversed sorted list
 - When $A[p]$ is inserted into the sorted $A[0..p-1]$, we need to compare $A[p]$ with all elements in $A[0 .. p-1]$ and move each element one position to the right
 $\Rightarrow \Omega(p)$ steps
 - the total number of steps is
$$\Omega(\sum_{i=1}^{N-1} i) = \Omega(N(N-1)/2) = \Omega(N^2)$$

Insertion Sort: Best-case Analysis

- The input is already sorted in increasing order
 - When inserting $a[p]$ into the sorted $a[0..p-1]$, only need to compare $a[p]$ with $a[p-1]$ and there is no data movement
 - For each iteration of the outer for-loop, the inner for-loop terminates after checking the loop condition once
=> $O(N)$ time
- If input is nearly sorted, insertion sort runs fast.

Insertion Sort: Summary

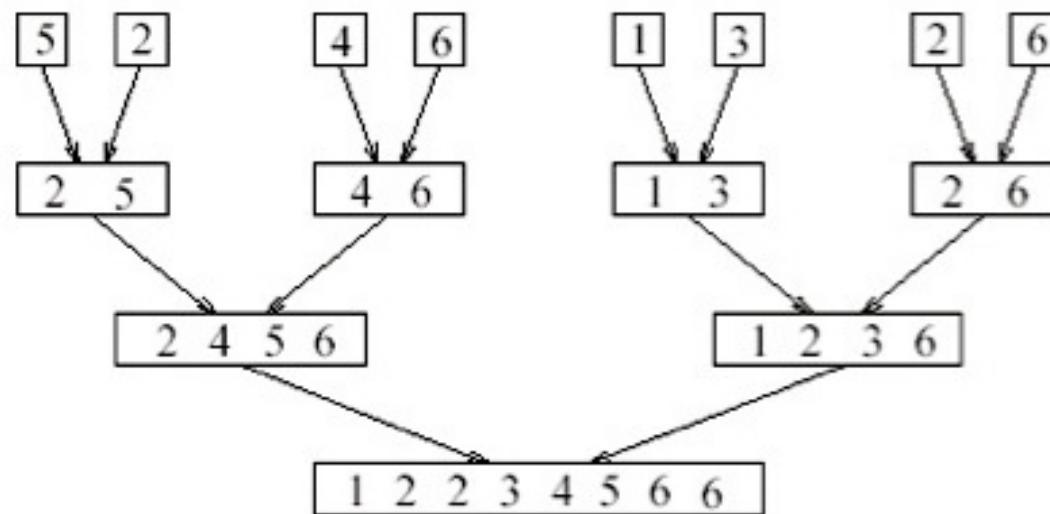
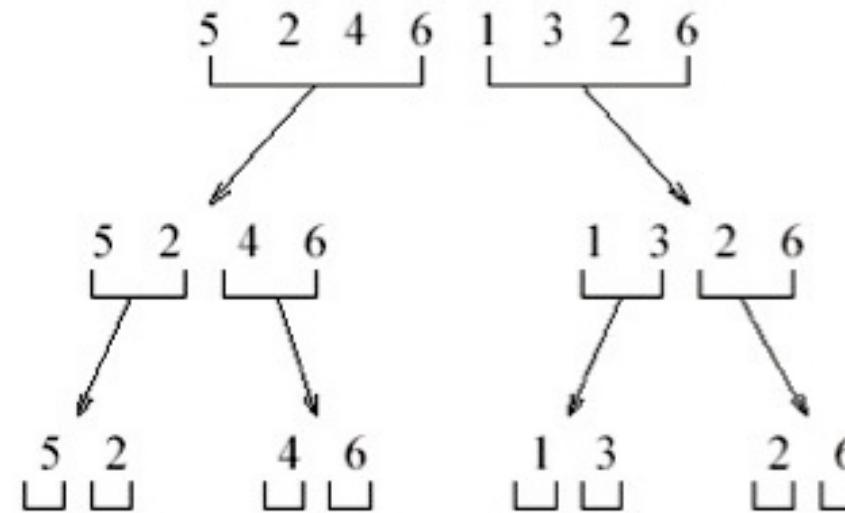
- Simple to implement.
- Efficient on (quite) small data sets.
- Efficient on data sets which are already substantially sorted.
- More efficient in practice than most other simple $O(N^2)$ algorithms such as selection sort or bubble sort.
 - it is linear in the best case.
- Properties:
 - stable
 - in-place
 - online algorithm

Comparison-based Sorting Algorithms by Divide-and-Conquer

Mergesort

- Procedure:
 1. Split data into roughly 2 equal halves.
 2. Recursively mergesort each half.
 3. Merge the two sorted halves together.
- 2 issues
 - #1: How to divide?
 - #2: How to merge?

Mergesort: Example



Mergesort: How to Divide?

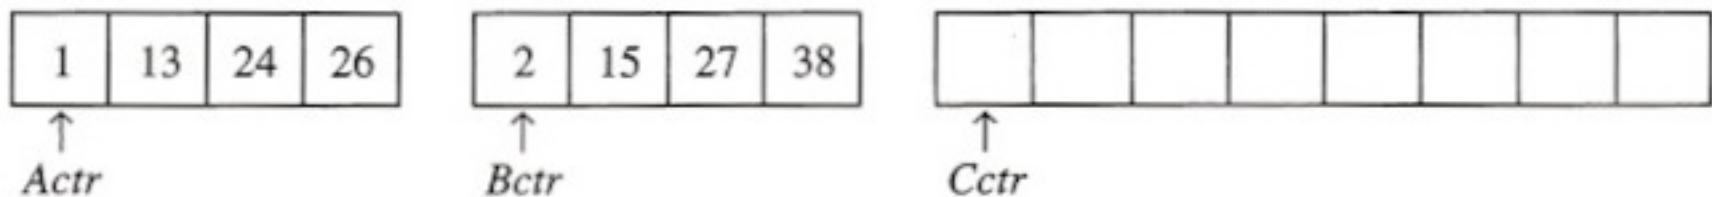
- Divide into 2 halves as equal as possible.

```
// Actual recursive routine requires an temporary array to store results
void mergesort(int* data, int* temp, int begin, int end)
{
    if (begin < end)
    {
        int center = (begin + end)/2;
        mergesort(data, temp, begin, center);
        mergesort(data, temp, center+1, end);
        merge(data, temp, begin, center, end);
    }
}

void mergesort(int* data, int begin, int end) // Driver function
{
    int* temp = new int [end - begin + 1]; // Store merged results
    mergesort(data, temp, begin, end);      // Actual recursive routine
    delete temp;
}
```

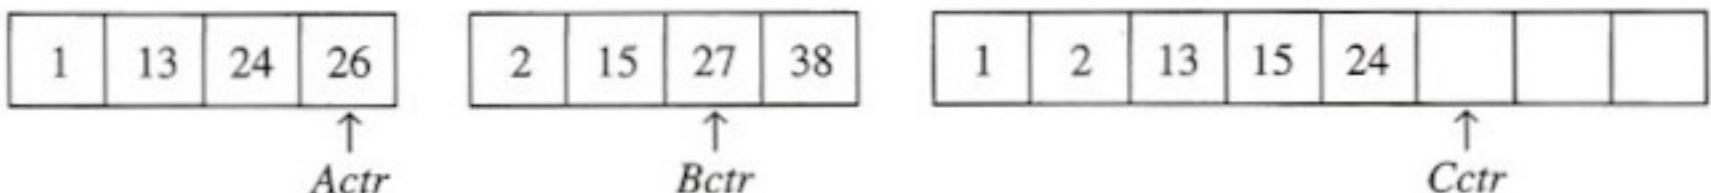
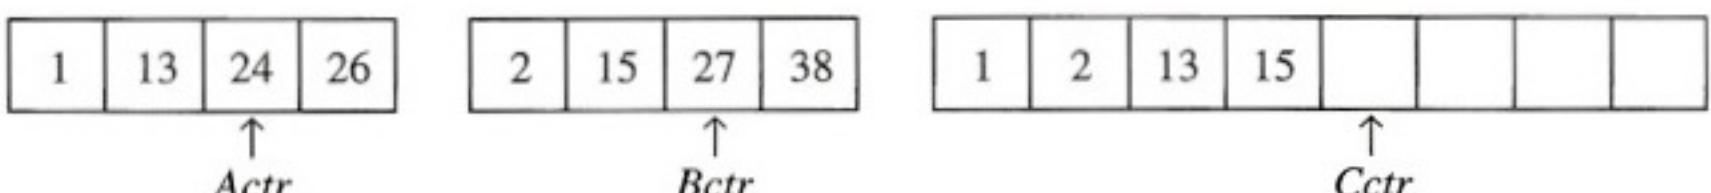
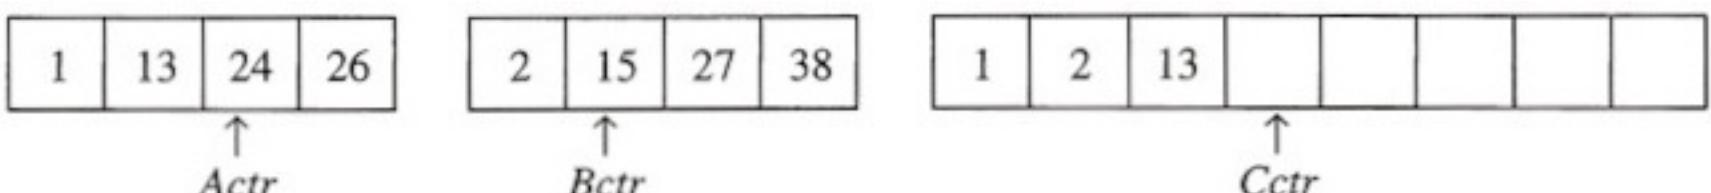
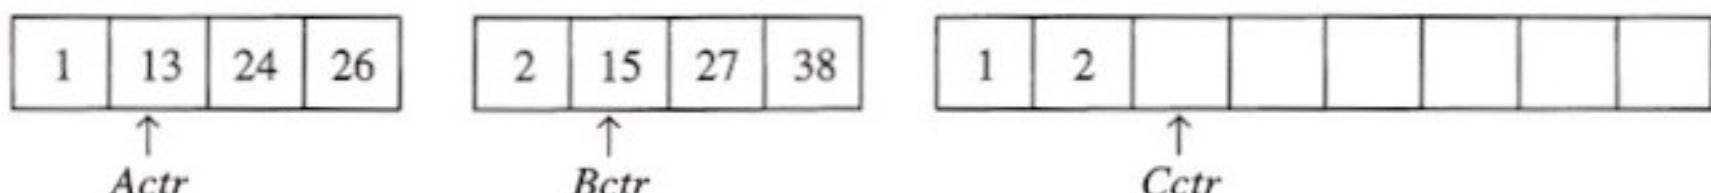
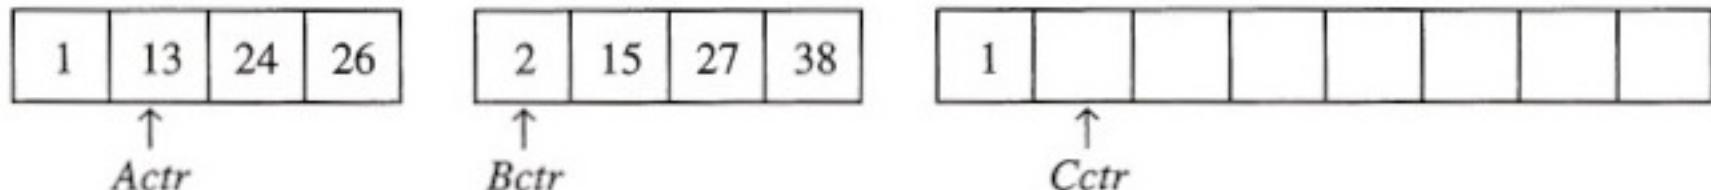
Mergesort: How to Merge?

- **Input:** 2 sorted arrays, A and B
- **Output:** an output sorted array C
- **Procedure:**
 1. 3 counters: Actr, Bctr, and Cctr, initially set to the beginning of their respective arrays.

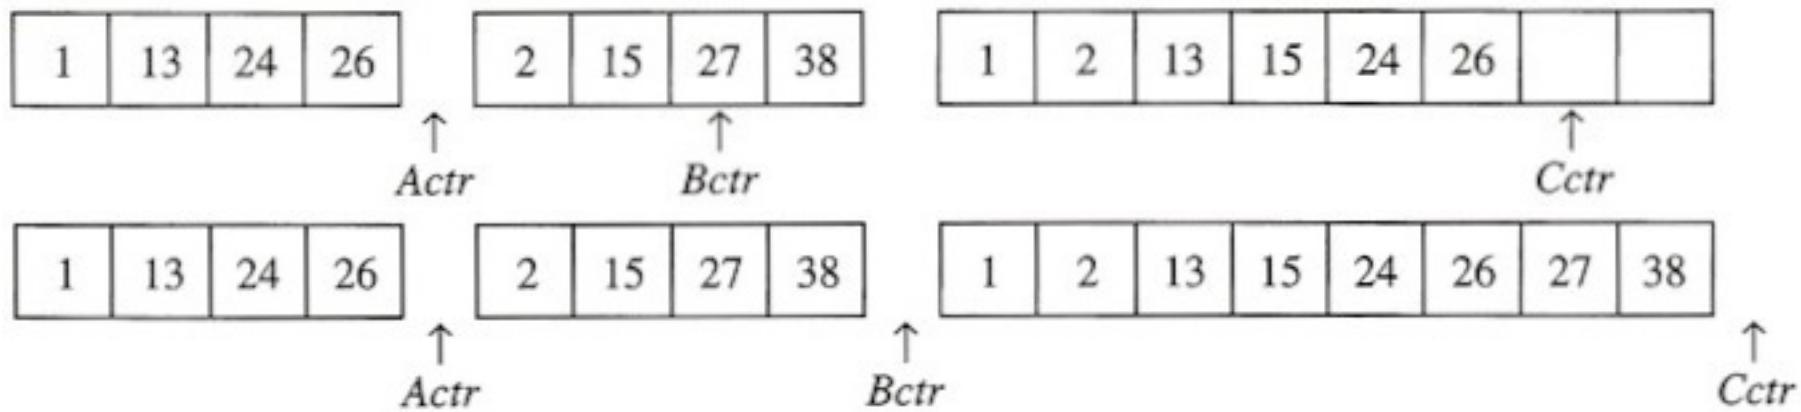


2. Compare A[Actr] and B[Bctr]: the smaller one is copied to the C[Cctr].
3. Advance Cctr and Actr/Bctr if its item is copied; repeat 2.
4. When either input list is exhausted, the remainder of the other list is copied to C.

Mergesort: How to Merge? ..



Mergesort: How to Merge? ...



- Running time analysis of merging 2 lists of lengths N_1 and N_2 :
 - $O(N_1 + N_2) = O(N)$ as $N = N_1 + N_2$.
- Space requirement:
 - need the space of an additional list

Merge Code

```
void merge(int* data, int* temp, int l_begin, int l_end, int r_end)
{
    int left = l_begin;      // Indexing the left subset
    int right = l_end + 1;   // Indexing the right subset
    int j = l_begin;         // Indexing the temp array for intermediate results

    // Merge the 2 sorted sublists
    while (left <= l_end && right <= r_end)
        temp[j++] = (data[left] <= data[right]) ? data[left++] : data[right++];

    // Copy any remaining items in the left sublist
    while (left <= l_end)
        temp[j++] = data[left++];

    // Copy any remaining items in the right sublist
    while (right <= r_end)
        temp[j++] = data[right++];

    // Copy results from the temp array back to the data array
    for (int k = l_begin; k <= r_end; ++k)
        data[k] = temp[k];
}
```

Mergesort: Analysis

- Let $T(N)$ be the worst-case running time of mergesort.
- Assume there are $N = 2^k$ numbers.
- Divide step: $O(1)$ to find the center; result in 2 \sim equal halves.
- Conquer step: $2 T(N/2)$
- Merge step: $O(N)$
- Recurrence equation:
 - $T(1) = 1$
 - $T(N) = 2T(N/2) + N$

Mergesort: Analysis ..

$$\begin{aligned}T(N) &= 2T\left(\frac{N}{2}\right) + N \\&= 2\left(2T\left(\frac{N}{4}\right) + \frac{N}{2}\right) + N \\&= 4T\left(\frac{N}{4}\right) + 2N \\&= 4\left(2T\left(\frac{N}{8}\right) + \frac{N}{4}\right) + 2N \\&= 8T\left(\frac{N}{8}\right) + 3N = \dots \\&= 2^k T\left(\frac{N}{2^k}\right) + kN\end{aligned}$$

Since $N = 2^k$, we have $k = \log_2 N$

$$\begin{aligned}T(N) &= 2^k T\left(\frac{N}{2^k}\right) + kN \\&= N + N \log N \\&= O(N \log N)\end{aligned}$$

Mergesort vs. Insertion Sort: An Experiment

- Code from textbook (using template)
- Unix time utility

n	Isort (secs)	Msort (secs)	Ratio
100	0.01	0.01	1
1000	0.18	0.01	18
2000	0.76	0.04	19
3000	1.67	0.05	33.4
4000	2.90	0.07	41
5000	4.66	0.09	52
6000	6.75	0.10	67.5
7000	9.39	0.14	67
8000	11.93	0.14	85

Mergesort vs. Insertion Sort: An Experiment ..

Comparing $n \log_{10} n$ and n^2

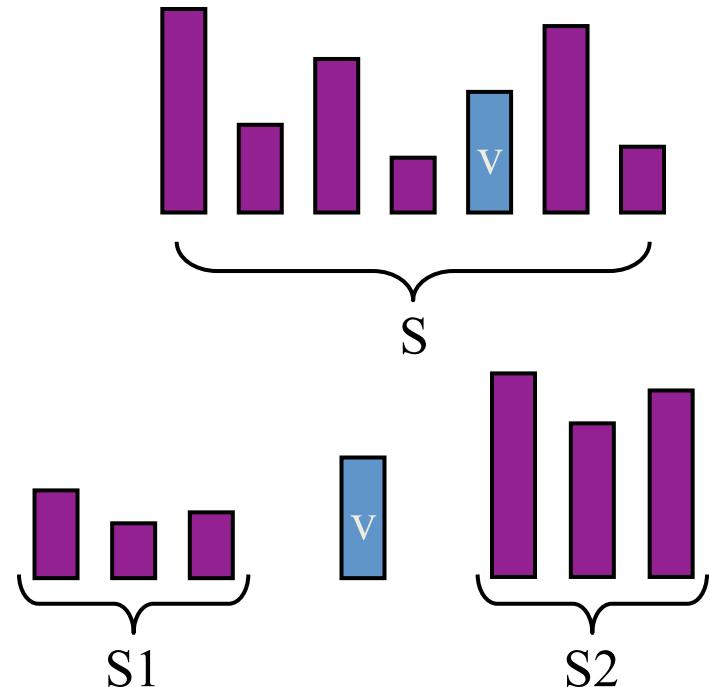
n	$n \log_{10} n$	n^2	Ratio
100	0.2K	10K	50
1000	3K	1M	333.33
2000	6.6K	4M	606
3000	10.4K	9M	863
4000	14.4K	16M	1110
5000	18.5K	25M	1352
6000	22.7K	36M	1588
7000	26.9K	49M	1820
8000	31.2K	64M	2050

Quicksort: Generally the Best

- **Fastest** known sorting algorithm in practice.
- **Best case:**
 - $O(N)$ [3-way partitioning (smaller, equal, larger subsets)
+ all values being equal]
 - $O(N \log N)$ [bad partition]
- **Average case:** $O(N \log N)$
- **Worst case:** $O(N^2)$
 - But, the worst case seldom happens.

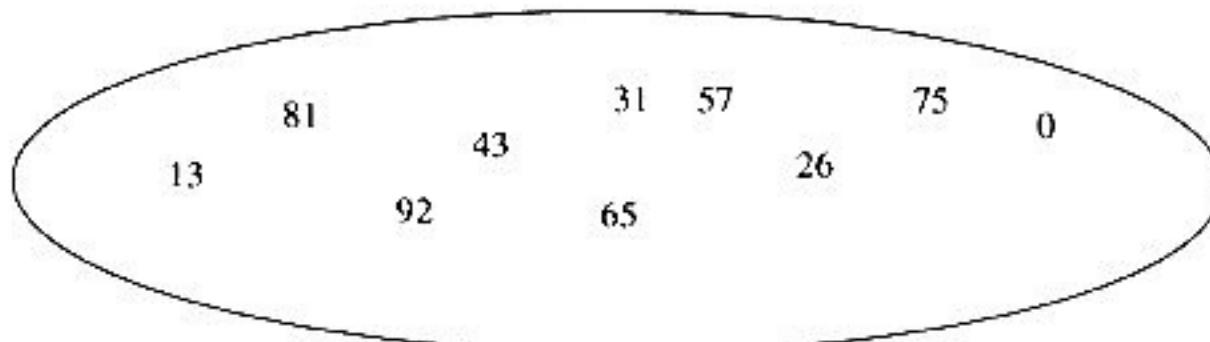
Quicksort: Idea

- Divide step:
 - Pick any element (**pivot**) v in S
 - Partition $S - \{v\}$ into 2 disjoint groups:
$$S_1 = \{x \in S - \{v\} \mid x \leq v\}$$
$$S_2 = \{x \in S - \{v\} \mid x \geq v\}$$
- Conquer step:
 - recursively sort S_1 and S_2
- Combine step:
 - just join together the sorted S_1 , followed by v , followed by the sorted S_2
(nothing extra needs to be done!!)

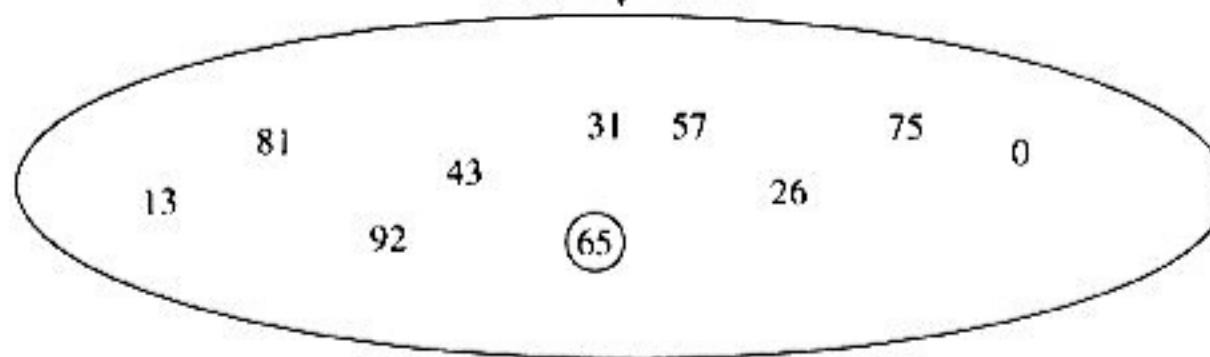


To simplify, we may assume that we don't have repetitive elements. So to ignore the 'equality' case.

Quicksort: Example

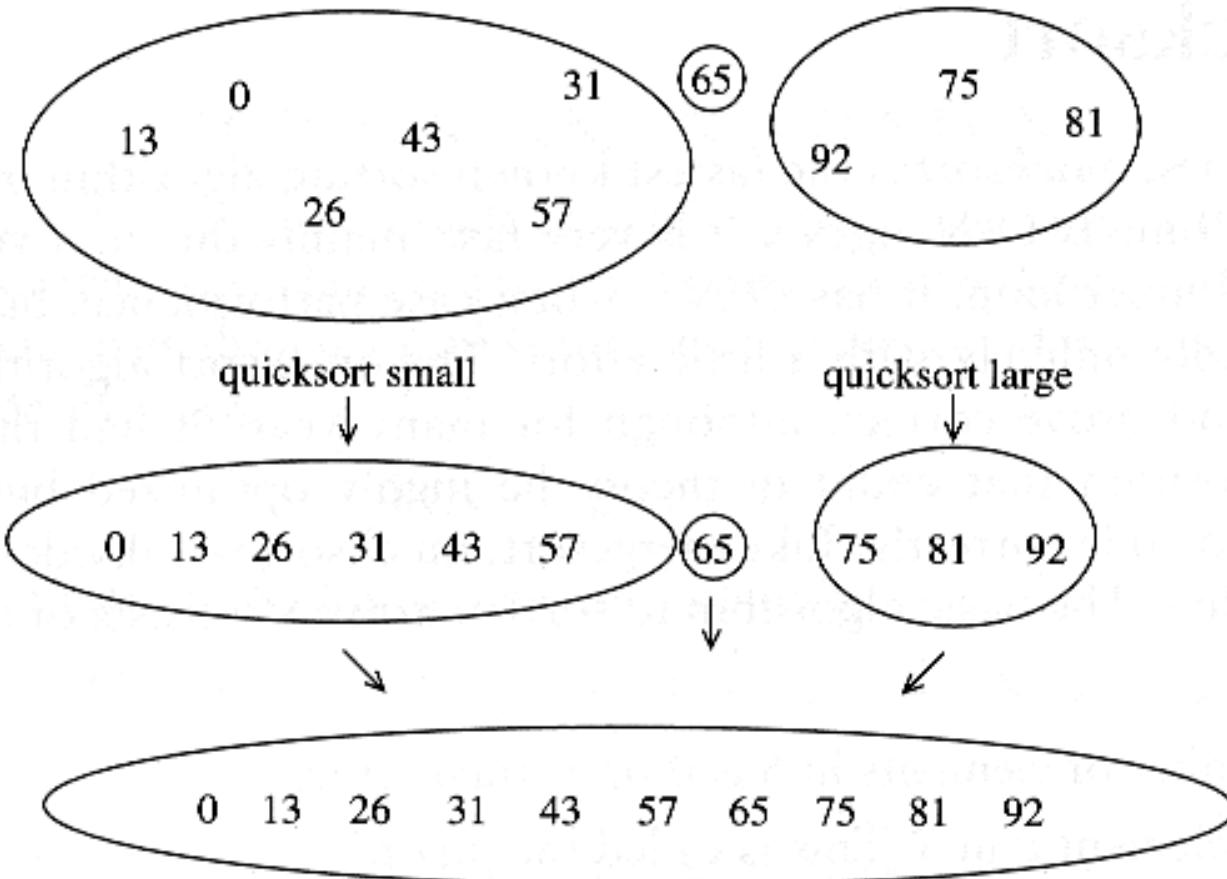


select pivot
↓



partition
↓

Quicksort: Example ..



Basic Quicksort Code

```
void quicksort(int* data, int begin, int end)
{
    int pivot_index = partition(data, begin, end); // Partition to 2 subsets
    quicksort(data, begin, pivot_index - 1);        // Sort the smaller items
    quicksort(data, pivot_index + 1, end);           // Sort the greater items
}
```

- Many ways to implement, but even the **slightest deviations** may cause surprisingly **bad results**.
 - it turns out it is hard to write it correctly ☹
 - many efficient implementation are not stable: the partition step does not preserve the ordering of elements with the same value.
- **Desirable**: sizes of the 2 subsets are as equal as possible
- **Major issue**: How to determine a good pivot?

Using the 1st Element as the Pivot?

- Use the **first element** as pivot
 - if the inputs are really **random** => **ok**
 - if the input is **pre-sorted** (or in reverse order)
 - all the elements go into S2 (or S1)
 - this happens consistently throughout the recursive calls
=> results in **O(n²)** behavior
- Choose the pivot **randomly**
 - generally safe
 - but random number generation itself can be **expensive**

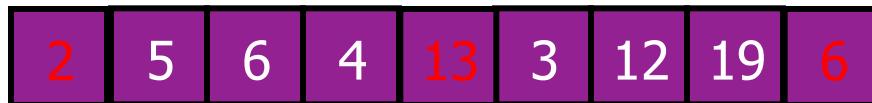
Exact Median as the Pivot?

- If the **median** of the data array is used
 - Partitioning always cuts the array into 2 equal subsets
- An optimal quicksort: $O(N \log N)$
- However, it is **hard** to find the **exact median**
 - chicken-egg problem
 - e.g., need to sort an array to pick the value in the middle!
- Need to find an approximation to the exact median.

Quicksort Pivot: Median of 3 --- Algorithm

- Compare just three elements: the **leftmost**, **rightmost** and the **center** elements.
- Swap these elements if necessary so that
 - $A[\text{begin}]$ = the smallest
 - $A[\text{center}]$ = median of the three elements
 - $A[\text{end}]$ = the largest
- Pick $A[\text{center}]$ as the pivot
- Swap $A[\text{center}]$ and $A[\text{end} - 1]$ so that pivot is at second last position (why?)

Quicksort Pivot: Median of 3 --- Example



$A[\text{begin}] = 2$,
 $A[\text{center}] = 13$, $A[\text{end}] = 6$



Swap $A[\text{center}]$ and $A[\text{end}]$



Choose $A[\text{center}]$ as **pivot**



Swap pivot and $A[\text{end}-1]$



Note we only need to partition $A[\text{begin}+1 \dots \text{end}-2]$. Why?

Quicksort Pivot: Median of 3 --- Code

```
#include <algorithm>
const int& median3(int* data, int begin, int end)
{
    int center = (begin + end)/2;

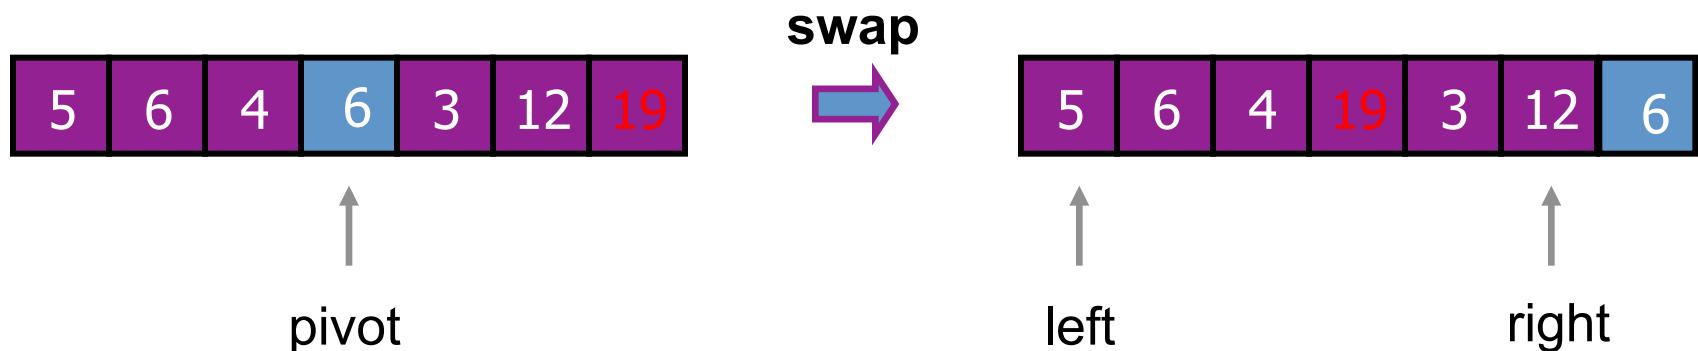
    if (data[begin] > data[center])
        std::swap(data[begin], data[center]);
    if (data[begin] > data[end])
        std::swap(data[begin], data[end]);

    // Now data[begin] is the smallest among the 3 items
    if (data[center] > data[end])
        std::swap(data[center], data[end]);

    // Place pivot as the last but one item
    std::swap(data[center], data[end - 1]);
    return data[end - 1];
}
```

Quicksort: Partitioning

- Want to partition an array $A[\text{begin} \dots \text{end}]$ in place.
- First, get the pivot element out of the way by swapping it with the last element. (Swap pivot and $A[\text{end}]$)
- Let left start at the first element and right start at the next-to-last element: $\text{left} = \text{begin}$, $\text{right} = \text{end}-1$

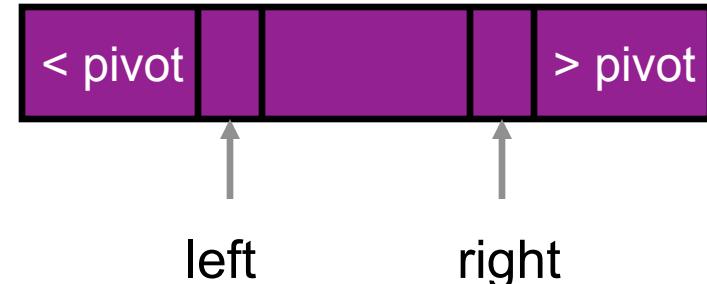


Quicksort: Partitioning ..

- Want to have

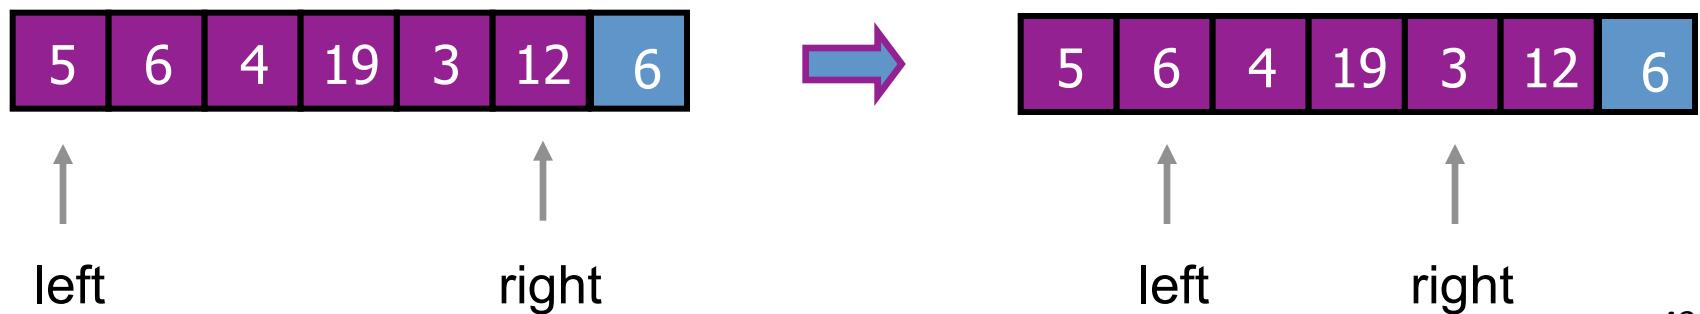
$A[x] \leq \text{pivot}$, for $x < \text{left}$

$A[x] \geq \text{pivot}$, for $x > \text{right}$



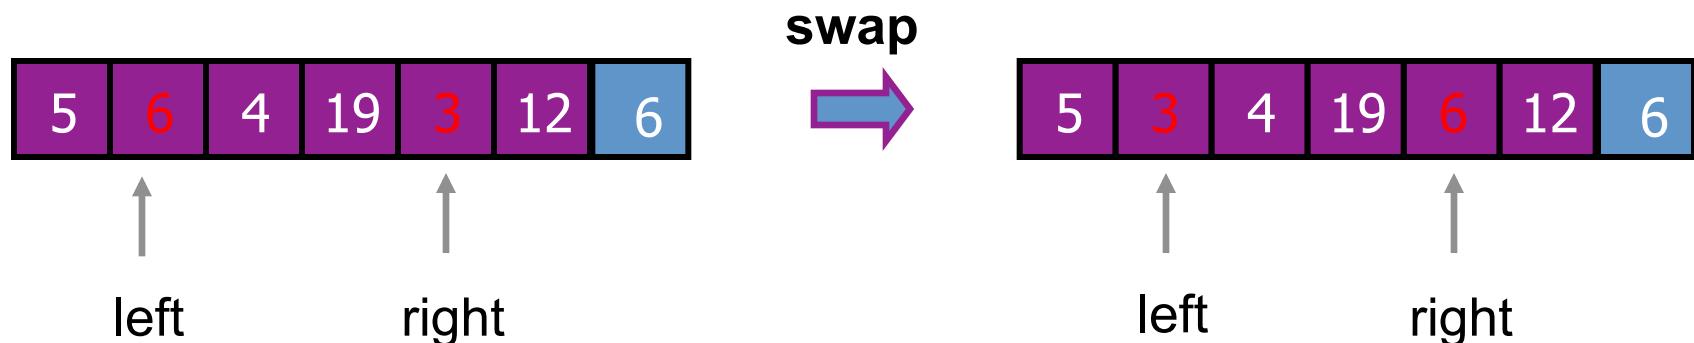
- As long as $\text{left} < \text{right}$

- Move left to the right, skipping over elements $<$ pivot
- Move right to the left, skipping over elements $>$ pivot
- When both left and right have stopped
 - $A[\text{left}] \geq \text{pivot}$ and $A[\text{right}] \leq \text{pivot}$



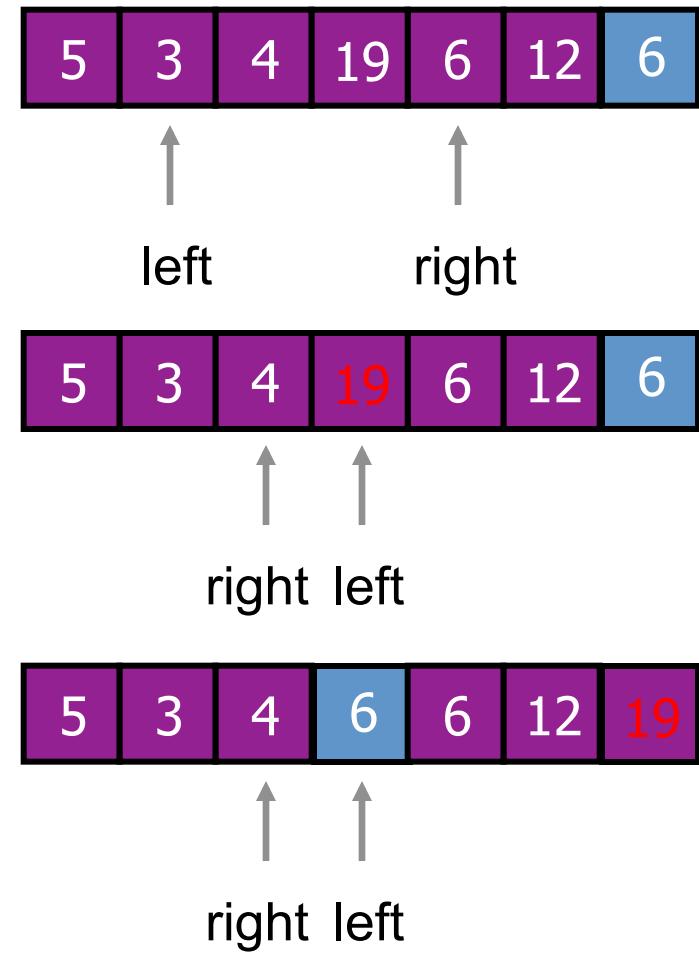
Quicksort: Partitioning ...

- When *left* and *right* stop, and *left* is to the left of *right*
 - Swap $A[\text{left}]$ and $A[\text{right}]$
 - larger element is pushed to the right and smaller element is pushed to the left
 - After swapping
 - $A[\text{left}] \leq \text{pivot}$ and $A[\text{right}] \geq \text{pivot}$
 - Repeat the process until *left* and *right* cross each other.



Quicksort: Partitioning

- When *left* and *right* cross
 - Swap $A[\text{left}]$ and pivot
- *right* will not run past the beginning
 - because $A[\text{begin}] \leq \text{pivot}$
- *left* will not run past the end
 - because $A[\text{end}-1] = \text{pivot}$
- Result:
 - $A[x] \leq \text{pivot}$, for $x < \text{left}$
 - $A[x] \geq \text{pivot}$, for $x > \text{right}$



Quicksort: Partitioning Code

```
int partition(int* data, int begin, int end) // Return the pivot index
{
    const int& pivot = median3(data, begin, end);
    int left = begin;
    int right = end - 1;

    while (true)
    {
        while (data[++left] < pivot) { } // Skip all smaller items
        while (data[--right] > pivot) { } // Skip all bigger items

        if (left < right)
            std::swap(data[left], data[right]);
        else
            break; // The 2 indices cross over
    }

    std::swap(data[left], data[end-1]); // Put pivot in its place
    return left;
}
```

Partitioning and Duplicates

- Partitioning so far defined is ambiguous for **duplicate** elements (the equality is included for both sets).
- Its '**randomness**' makes a '**balanced**' distribution of duplicate elements in the 2 subsets.
- When all elements are identical:
 - both left and right stop right away in each iteration
=> many swaps
 - but they cross in the middle
=> partition is balanced
=> $O(N \log N)$

Quicksort is not Good for Short arrays

- For very small arrays, quicksort does not perform as well as insertion sort.
 - How small depends on many factors, such as the time spent making a recursive call, the compiler, etc.
- Do **not** use quicksort recursively for small arrays.
 - Instead, use a sorting algorithm that is efficient for small arrays, such as **insertion sort**.

Practical Implementation of Quicksort

```
void quicksort(int* data, int begin, int end)
{
    int insertion_sort(int* data, int N); // Forward declaration

    // >= 2 for median-3 partitioning, but 10-20 for efficiency
    const int QSORT_MIN_SIZE = 10;

    if (end - begin > QSORT_MIN_SIZE)
    {
        // Partition to 2 subsets of similar sizes
        int pivot_index = partition(data, begin, end);

        // Recursively sort the smaller items
        quicksort(data, begin, pivot_index - 1);

        // Recursively sort the greater items
        quicksort(data, pivot_index + 1, end);
    }
    else // Insertion sort is better for short list of data
        insertion_sort(data + begin, end - begin + 1);
}
```

Quicksort Analysis

- Assumptions
 - A **random** pivot (no median-of-three partitioning)
 - No cutoff for small arrays
- Running time
 - **pivot selection**: constant time = $O(1)$
 - **partitioning**: linear time = $O(N)$
 - running time of the 2 recursive calls
- $T(N) = T(M) + T(N-M-1) + cN$
 - c is a constant
 - M is the number of elements in S_1

Worst-Case Analysis

- What will be the worst case?
 - The **pivot** is the **smallest** element, all the time
 - **Partition** is always **unbalanced**

$$T(N) = T(N - 1) + cN$$

$$T(N - 1) = T(N - 2) + c(N - 1)$$

$$T(N - 2) = T(N - 3) + c(N - 2)$$

⋮

$$T(2) = T(1) + c(2)$$

$$T(N) = T(1) + c \sum_{i=2}^N i = O(N^2)$$

Best-case Analysis

- What will be the best case?
 - Partition is perfectly balanced.
 - Pivot is always in the middle (median of the array)

$$T(N) = 2T(N/2) + cN$$

$$\frac{T(N)}{N} = \frac{T(N/2)}{N/2} + c$$

divide the whole equation by N

$$\frac{T(N/2)}{N/2} = \frac{T(N/4)}{N/4} + c$$

$$\frac{T(N/4)}{N/4} = \frac{T(N/8)}{N/8} + c$$

⋮

$$\frac{T(2)}{2} = \frac{T(1)}{1} + c$$

$$\frac{T(N)}{N} = \frac{T(1)}{1} + c \log N$$

add up all equations; remove identical terms on both sides

$$T(N) = N + cN \log N = O(N \log N)$$

Average-Case Analysis

- Assume
 - Each of the sizes for S1 is equally likely
- This assumption is valid for our pivoting (median-of-three) strategy
- On average, the running time is $O(N \log N)$

Quicksort vs. Mergesort

- Both **quicksort** and **mergesort** take $O(N \log N)$ in the average case.
- **Quicksort**: the inner loop consists of an increment/decrement (by 1, which is fast), a test and a jump.
- **Partitioning** in **quicksort** and **merging** in **mergesort** both take $O(N)$ time.
- But, **mergesort** requires **copying** results back to the original data array => may be slower.
- Exact performance depends on the speed of data copying, data comparison, etc.
 - Java sorting lib: **mergesort**
 - C++ STL sort: **quicksort**

Lower Bound for Sorting Linear-Time Sorting Algorithms

Lower Bound for Sorting

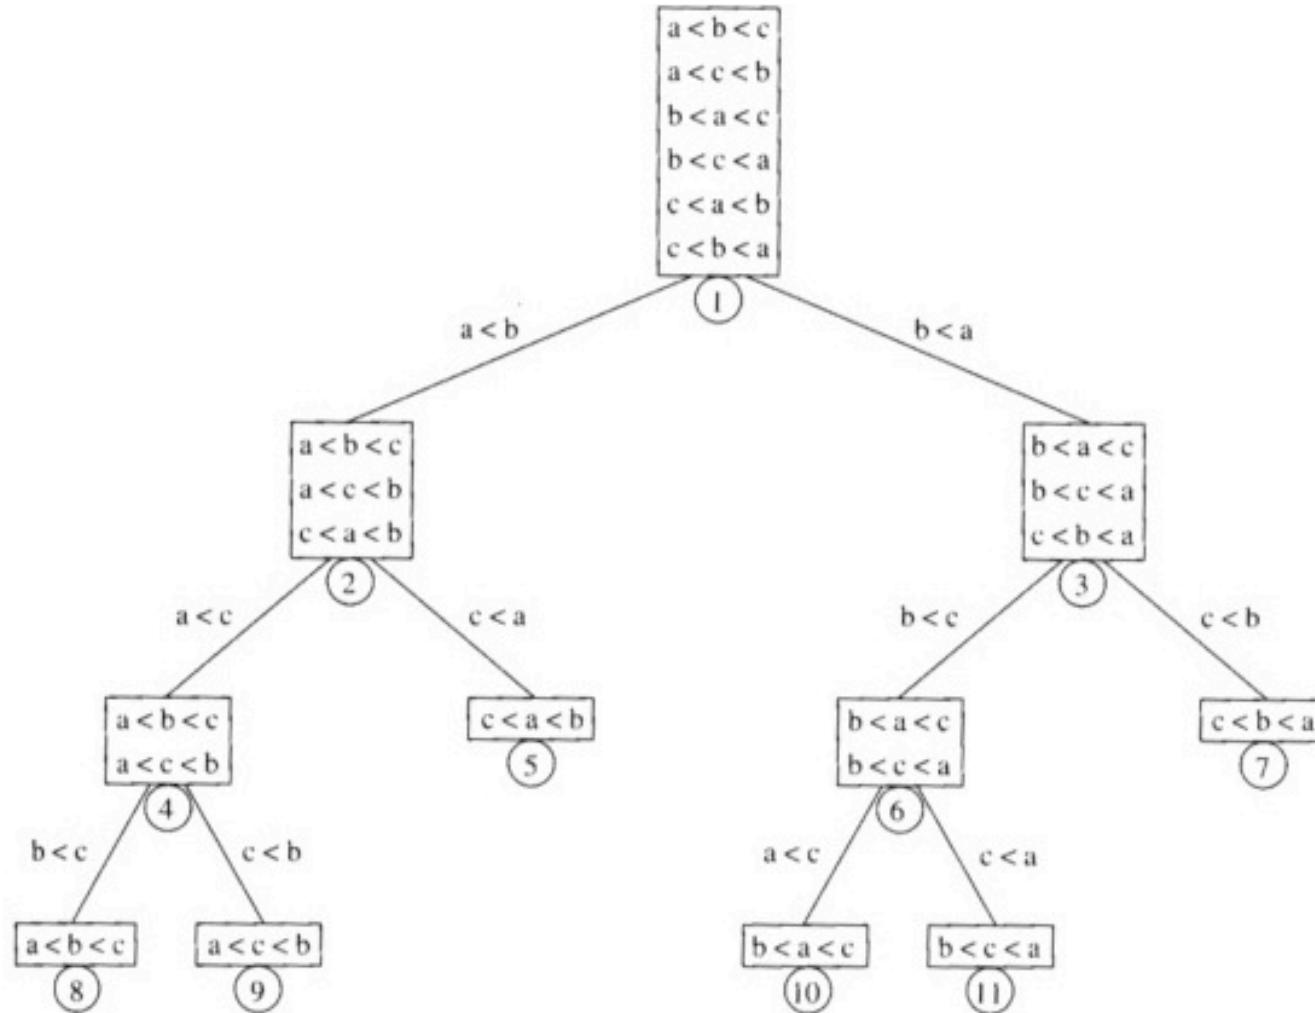
- Mergesort
 - worst-case running time is $O(N \log N)$
- Are there better algorithms?
- Theorem:

Any sorting algorithm based only on comparisons to sort N elements takes $\Omega(N \log N)$ comparisons in the worst case (worse-case input).

Decision Trees for Sorting

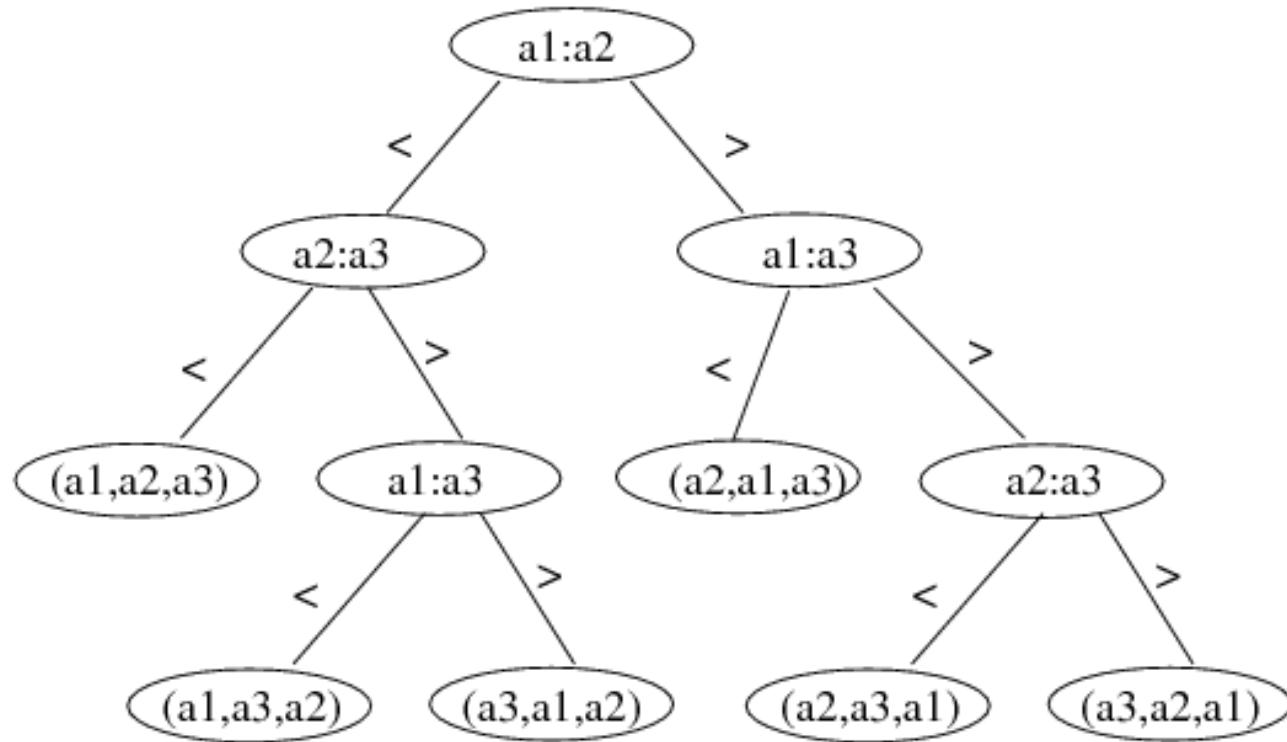
- Suppose we want to sort N distinct elements.
- How many **possible orderings** do we have for N elements?
 - There are **$N!$** possible orderings
 - e.g., the sorted output for { a, b, c } can be one of
 - a b c, b a c, a c b, c a b, c b a, b c a.
- Any comparison-based sorting process can be represented as a binary **decision tree**.
 - Each **node** represents a set of possible orderings, consistent with all the comparisons that have been made.
 - Tree **edges** are results of intermediate comparisons.
 - Leaves are the results.

Decision Tree for Sorting 3 Elements



Decision Tree for Insertion Sort

- Different algorithms have different decision trees.
- There exists a **root-to-leaf path** that arrives at a leaf containing inputs in **sorted order**.
- For **insertion sort**, the longest path in such a tree is **$O(N^2)$** .



Lower Bound for Sorting

- Worst-case: number of comparisons used by a sorting algorithm is equal to the depth of the deepest leaf.
- Average case: number of comparisons is equal to the average depth of the leaves
- A decision tree to sort N elements must have $N!$ leaves.
 - a binary tree of depth d has at most 2^d leaves
⇒ a binary tree with 2^d leaves must have depth at least d
⇒ the decision tree with $N!$ leaves must have depth at least $\lceil \log_2(N!) \rceil$
- Thus, any sorting algorithm based only on comparisons between elements requires at least $\lceil \log_2(N!) \rceil$ comparisons in the worst case.

Lower Bound for Sorting ..

$$\begin{aligned}\log(N!) &= \log(N(N-1)(N-2)\cdots 2 \cdot 1) \\&= \log N + \log(N-1) + \log(N-2) + \cdots + \log 2 + \log 1 \\&\geq \log N + \log(N-1) + \log(N-2) + \cdots + \log(N/2) \\&\geq \frac{N}{2} \log \frac{N}{2} \\&= \frac{N}{2} \log N - \frac{N}{2} \\&= \Omega(N \log N)\end{aligned}$$

- Therefore, any sorting algorithm based on comparisons between elements requires $\Omega(N \log N)$ comparisons.

Linear Time Sorting Algorithms

- Can we do better?
- Yes, we may even obtain $O(N)$ sorting algorithm if the input has special structure or extra information, e.g.,
 - uniformly distributed
 - every number can be represented by r digits
- Linear time sorting algorithms
 - bucket sort
 - radix sort

Bucket Sort with No Duplicates

- To sort N +ve integers $A[1 .. N]$ in the range of 1 to $M \geq N$.
These integers may be keys of actual storing objects.
- 1. Define an array $B[1 .. M]$; initialize all to 0 $\Rightarrow O(M)$
- 2. Scan through the input list $A[i]$; insert $A[i]$ into $B[A[i]]$ $\Rightarrow O(N)$
- 3. Scan B once and read out the non-zero integers $\Rightarrow O(M)$
- Total time: $O(M + N)$
 - if M is $O(N)$, then total time is $O(N)$.
 - Can be bad if range is very big, e.g. $M = O(N^2)$.
- Example: To sort { 8 1 9 5 2 6 3 }. $N = 7$, $M = 9$.



- Output: 1 2 3 5 6 8 9

Bucket Sort with Duplicates

- B is an array of **linked lists**.
- Each cell in B has 2 pointers:
 - head points to the beginning of a linked list.
 - tail points to the end of the linked list.
- $A[j]$ is inserted at the **end** of the list headed by $B[A[j]]$
=> **stable sort**
- After all elements of array A are inserted, array B is sequentially traversed and each non-empty list is printed out **from head to tail**.
- Time: **O(M + N)**

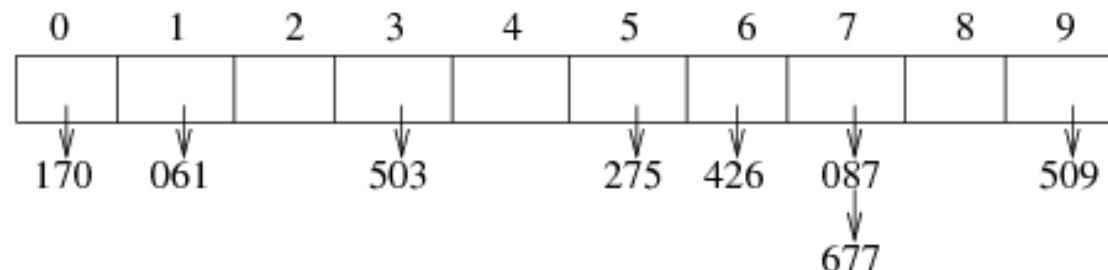
Radix Sort

- Given N integers in the range 0 to $(r^k - 1)$
 - i.e., they can be represented by at most k digits
 - $d_1d_2\dots d_k$ where d_i are digits in base r
 - d_1 : most significant digit
 - d_k : least significant digit
- Algorithm
 1. Bucket sort by the least significant digit first.
=> numbers with the same d_k digit go to same bin
 2. Re-order the numbers according to their current sorted order from bucket 0 to bucket $(r-1)$.
 3. Sort by the next least significant digit.
 4. Continue this process until the numbers are sorted on all k digits.

Radix Sort: Example 1

- Example: 275, 087, 426, 061, 509, 170, 677, 503
- 1st pass: Least-significant-digit first.

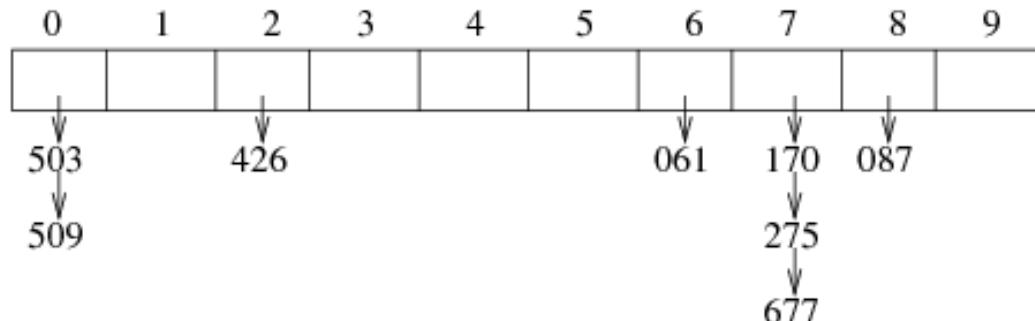
1st pass



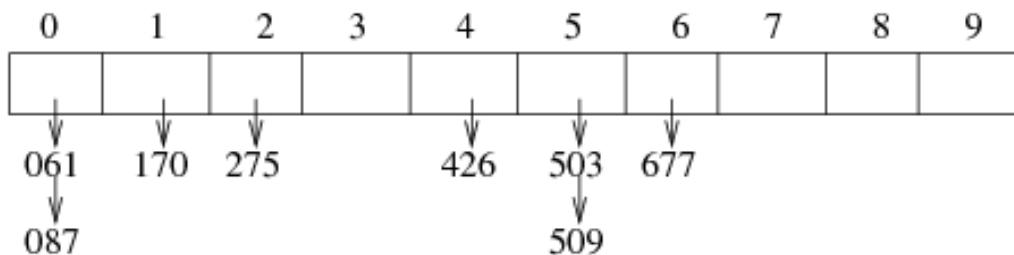
- Read the currently sorted numbers from the buckets:
 - 170 061 503 275 426 087 677 509

Radix Sort: Example 1 ..

2nd pass Before: 170 061 503 275 426 087 677 509



3rd pass Before: 503 509 426 061 170 275 677 087



in sorted order

After: 061 087 170 275 426 503 509 677

Radix Sort: Example 2

- One commonly uses **radix sort** to sort cards.
- Use 2 digits to represent each card: d_1d_2
 - $d_1 = \spadesuit \heartsuit \clubsuit \diamondsuit$: base 4
 - $\spadesuit \leq \clubsuit \leq \heartsuit \leq \diamondsuit$
 - $d_2 = A, 2, 3, \dots J, Q, K$: base 13
 - $A \leq 2 \leq 3 \leq \dots \leq J \leq Q \leq K$
 - e.g., $\spadesuit 2 \leq \clubsuit 2 \leq \heartsuit 5 \leq \diamondsuit K$

Radix Sort: Pseudo-Code

- $A = \text{input array}$, $n = \text{number of data}$, $d = \# \text{ of digits}$

Algorithm $\text{RadixSort}(A, n, d)$

```
1.   for  $0 \leq p \leq 9$            ← assume base = 10
2.     do  $Q[p] := \text{empty queue};$ 
3.    $D := 1;$ 
4.   for  $1 \leq k \leq d$            ← Consider the k-th digit
5.     do
6.        $D := 10 * D;$ 
7.       for  $0 \leq i < n$            ← Scan A[i], put into correct bucket
8.         do  $t := (A[i] \bmod D) \div (D/10);$ 
9.           enqueue( $A[i]$ ,  $Q[t]$ );
10.       $j := 0;$ 
11.      for  $0 \leq p \leq 9$            ← Put numbers back to the original array
12.        do while  $Q[p]$  is not empty
13.          do  $A[j] := \text{dequeue}(Q[p]);$ 
14.           $j := j + 1;$ 
```

Radix Sort: Summary

- Increasing the **base** r decreases the number of digits k .
- Running time analysis:
 - k passes over the numbers: 1 bucket sort for each digit
 - each pass takes $2N$
 - total: $O(2Nk) = O(Nk)$
- Radix sort is **not** based on comparisons; the values are used as array indices.
- If all N input values are **distinct**, then $k = \Omega(\log N)$.
 - e.g., in binary digits, to represent 8 different numbers, we need at least 3 digits.
 - The running time of radix sort also becomes $\Omega(N \log N)$.